

# Data Collection and Labeling

CS 203: Software Tools and Techniques for AI

Prof. Nipun Batra

## Table of contents

<b>1</b>	<b>Module Overview</b>	<b>4</b>
1.1	Four Core Components . . . . .	4
<b>2</b>	<b>Part 1: Data Collection</b>	<b>4</b>
2.1	Why Data Collection Matters . . . . .	4
2.2	Common Data Sources . . . . .	5
2.2.1	Digital Sources . . . . .	5
2.2.2	Physical Sources . . . . .	5
2.3	Instrumentation: The Foundation . . . . .	5
2.3.1	Key Principles . . . . .	5
2.4	Basic Instrumentation Example . . . . .	5
2.5	Web Analytics Tools . . . . .	6
2.5.1	Google Analytics 4 . . . . .	6
2.5.2	Mixpanel . . . . .	6
2.6	Self-Hosted Analytics . . . . .	6
2.6.1	Why Self-Host? . . . . .	6
2.7	Plausible & Umami Examples . . . . .	7
2.8	Application Performance Monitoring . . . . .	7
2.8.1	Sentry . . . . .	7
2.8.2	Datadog . . . . .	7
2.9	Database Event Logging . . . . .	8
2.10	Change Data Capture (CDC) . . . . .	8
2.10.1	Debezium Example . . . . .	8
2.11	OpenTelemetry: Industry Standard . . . . .	9
2.12	Web Scraping with Scrapy . . . . .	9
2.13	Ethical Web Scraping . . . . .	10
2.14	Streaming Data with Kafka . . . . .	10
2.15	Data Collection Best Practices . . . . .	11
2.15.1	Schema Design with Pydantic . . . . .	11

2.16 Sampling Strategies . . . . .	11
2.17 Privacy and Compliance . . . . .	11
<b>3 Part 2: Data Validation . . . . .</b>	<b>12</b>
3.1 Why Data Validation Matters . . . . .	12
3.1.1 The Cost of Bad Data . . . . .	12
3.1.2 Real-World Impact . . . . .	12
3.2 Types of Data Quality Issues . . . . .	12
3.3 Data Validation with Pydantic . . . . .	13
3.4 Using Pydantic . . . . .	13
3.5 Great Expectations Framework . . . . .	14
3.6 Pandera: Statistical Validation . . . . .	14
3.7 Custom Validation Pipelines . . . . .	15
3.8 Statistical Distribution Validation . . . . .	15
3.9 Data Quality Monitoring . . . . .	16
<b>4 Part 3: Data Labeling . . . . .</b>	<b>17</b>
4.1 What is Data Labeling? . . . . .	17
4.1.1 Why It's Critical . . . . .	17
4.2 Types of Labeling Tasks . . . . .	17
4.2.1 Classification . . . . .	17
4.2.2 Object Detection . . . . .	17
4.2.3 Sequence Labeling . . . . .	17
4.2.4 Structured Prediction . . . . .	18
4.3 Label Studio . . . . .	18
4.3.1 Key Features . . . . .	18
4.3.2 Installation . . . . .	18
4.4 Label Studio: Image Classification . . . . .	18
4.5 Label Studio: Object Detection . . . . .	19
4.6 Label Studio: NER . . . . .	19
4.7 Alternative Labeling Tools . . . . .	19
4.7.1 Labelbox . . . . .	19
4.7.2 CVAT . . . . .	19
4.7.3 Prodigy . . . . .	20
4.7.4 Scale AI / SageMaker . . . . .	20
4.8 Inter-Annotator Agreement . . . . .	20
4.8.1 Cohen's Kappa ( ) . . . . .	20
4.9 Cohen's Kappa Implementation . . . . .	20
4.10 Fleiss' Kappa: Multiple Annotators . . . . .	21
4.11 Improving Agreement . . . . .	21
4.11.1 1. Clear Guidelines . . . . .	21
4.11.2 2. Training Phase . . . . .	22
4.11.3 3. Regular Consensus Meetings . . . . .	22

4.12	Sampling Strategies . . . . .	22
4.12.1	Random Sampling . . . . .	22
4.12.2	Stratified Sampling . . . . .	22
4.13	Snowball Sampling . . . . .	22
4.14	Active Learning . . . . .	23
4.14.1	Process . . . . .	23
4.15	Uncertainty Sampling . . . . .	23
4.16	Active Learning Loop . . . . .	24
4.17	Active Learning Benefits . . . . .	24
4.17.1	Cost Reduction . . . . .	24
4.17.2	Real-world Savings . . . . .	25
4.18	Weak Supervision with Snorkel . . . . .	25
4.19	Combining Weak Labels . . . . .	25
4.20	Pre-labeling with Models . . . . .	26
<b>5</b>	<b>Part 4: Data Augmentation</b>	<b>26</b>
5.1	Why Data Augmentation? . . . . .	27
5.1.1	The Problem . . . . .	27
5.1.2	The Solution . . . . .	27
5.2	Image Data Augmentation . . . . .	27
5.3	Advanced: Albumentations . . . . .	28
5.4	Cutout, Mixup, CutMix . . . . .	28
5.4.1	Cutout . . . . .	28
5.4.2	Mixup . . . . .	29
5.4.3	CutMix . . . . .	29
5.5	Text Data Augmentation . . . . .	29
5.5.1	Synonym Replacement . . . . .	29
5.6	EDA: Easy Data Augmentation . . . . .	30
5.7	Back-Translation . . . . .	30
5.8	Contextual Augmentation with BERT . . . . .	31
5.9	Audio Augmentation . . . . .	31
5.10	Time Series Augmentation . . . . .	32
5.11	SMOTE for Tabular Data . . . . .	32
5.12	SMOTE Variants . . . . .	33
5.12.1	ADASYN . . . . .	33
5.12.2	BorderlineSMOTE . . . . .	33
5.13	Generative Models . . . . .	33
5.14	Best Practices . . . . .	34
5.14.1	1. Domain-Appropriate . . . . .	34
5.14.2	2. Preserve Semantics . . . . .	34
5.14.3	3. Validation Split First . . . . .	34
5.15	Monitor Augmentation Impact . . . . .	34
5.16	Class-Specific Augmentation . . . . .	35

<b>6 Summary</b>	<b>36</b>
6.1 Key Takeaways . . . . .	36
6.1.1 Data Collection . . . . .	36
6.1.2 Data Validation . . . . .	36
6.2 Key Takeaways (cont.) . . . . .	36
6.2.1 Data Labeling . . . . .	36
6.2.2 Data Augmentation . . . . .	36
6.3 Hands-On Lab . . . . .	37
6.3.1 Complete Data Pipeline . . . . .	37
6.4 Additional Resources . . . . .	37
6.4.1 Documentation . . . . .	37
6.4.2 Key Papers . . . . .	37
<b>7 Questions?</b>	<b>37</b>
7.1 Thank You! . . . . .	37

## 1 Module Overview

### 1.1 Four Core Components

1. **Data Collection** - Tools and techniques for gathering data from diverse sources
2. **Data Validation** - Ensuring data quality and reliability
3. **Data Labeling** - Annotating datasets with ground truth
4. **Data Augmentation** - Expanding datasets strategically



Quality data is the foundation of successful AI systems. Garbage in, garbage out!

## 2 Part 1: Data Collection

Instrumenting and Logging Data Sources

### 2.1 Why Data Collection Matters

- **Real-world AI systems** depend on continuous data flow
- **User behavior** changes over time → models need fresh data
- **Production systems** require automated collection pipelines
- **Debugging** often requires understanding what data was seen

## 2.2 Common Data Sources

### 2.2.1 Digital Sources

- Web applications
- Mobile apps
- IoT devices
- APIs and databases

### 2.2.2 Physical Sources

- Sensors
- Cameras
- Microphones
- Manual entry

## 2.3 Instrumentation: The Foundation

**Instrumentation** = Adding code to collect data about system behavior

### 2.3.1 Key Principles

1. **Minimal Performance Impact** - Don't slow down production systems
2. **Comprehensive Coverage** - Capture all relevant events
3. **Privacy-Aware** - Respect user privacy and regulations (GDPR, CCPA)
4. **Structured Logging** - Consistent formats for easy parsing

## 2.4 Basic Instrumentation Example

```
import logging
import json
from datetime import datetime

def log_user_action(user_id, action, metadata):
    event = {
        "timestamp": datetime.utcnow().isoformat(),
        "user_id": user_id,
        "action": action,
        "metadata": metadata
```

```
    }
    logging.info(json.dumps(event))
```

## 2.5 Web Analytics Tools

### 2.5.1 Google Analytics 4

```
// Track custom events
gtag('event', 'search', {
  'search_term': query,
  'results_count': results.length
});
```

**Pros:** Free, Rich ecosystem, Real-time dashboards **Cons:** Privacy concerns, Limited customization

### 2.5.2 Mixpanel

```
// Track with properties
mixpanel.track('Video Played', {
  'video_id': video.id,
  'duration': video.length,
  'quality': '1080p'
});
```

**Pros:** Detailed analytics, Cohort analysis **Cons:** Costly at scale

## 2.6 Self-Hosted Analytics

### 2.6.1 Why Self-Host?

- **Data ownership** - Complete control over your data
- **Privacy compliance** - GDPR-friendly by design
- **No tracking cookies** - Lightweight and fast
- **Cost-effective** - Predictable infrastructure costs

## 2.7 Plausible & Umami Examples

```
// Plausible: Simple, privacy-focused
<script defer data-domain="yourdomain.com"
    src="https://plausible.io/js/script.js"></script>

plausible('signup', {props: {plan: 'premium'}});
```

```
// Umami: Open-source alternative
<script async src="https://analytics.yourdomain.com/umami.js"
    data-website-id="your-website-id"></script>

umami.track('button-click', {button: 'subscribe'});
```

## 2.8 Application Performance Monitoring

### 2.8.1 Sentry

```
import sentry_sdk

sentry_sdk.init(
    dsn="your-dsn",
    traces_sample_rate=1.0
)

try:
    process_data()
except Exception as e:
    sentry_sdk.capture_exception(e)
```

### 2.8.2 Datadog

```
from datadog import initialize, statsd

initialize(
    api_key='your-key',
    app_key='your-app-key'
```

```
)  
  
statsd.increment('api.requests')  
statsd.histogram('api.latency', 245)
```

## 2.9 Database Event Logging

```
-- PostgreSQL: Track all changes to users table  
CREATE TABLE user_audit (  
    audit_id SERIAL PRIMARY KEY,  
    user_id INT,  
    action VARCHAR(10),  
    old_data JSONB,  
    new_data JSONB,  
    changed_at TIMESTAMP DEFAULT NOW()  
);  
  
CREATE OR REPLACE FUNCTION audit_user_changes()  
RETURNS TRIGGER AS $$  
BEGIN  
    IF (TG_OP = 'UPDATE') THEN  
        INSERT INTO user_audit(user_id, action, old_data, new_data)  
        VALUES (NEW.id, 'UPDATE', row_to_json(OLD), row_to_json(NEW));  
    END IF;  
    RETURN NEW;  
END;  
$$ LANGUAGE plpgsql;
```

## 2.10 Change Data Capture (CDC)

**CDC** = Capturing and streaming database changes in real-time

### 2.10.1 Debezium Example

```
{  
    "name": "inventory-connector",  
    "config": {
```

```

    "connector.class": "io.debezium.connector.postgresql.PostgresConnector",
    "database.hostname": "postgres",
    "database.port": "5432",
    "table.include.list": "public.orders,public.customers",
    "topic.prefix": "dbserver1"
}
}

```

**Benefits:** Real-time streaming, No app changes, Complete history

## 2.11 OpenTelemetry: Industry Standard

```

from opentelemetry import trace
from opentelemetry.sdk.trace import TracerProvider

trace.set_tracer_provider(TracerProvider())
tracer = trace.get_tracer(__name__)

@tracer.start_as_current_span("process_request")
def process_request(user_id):
    span = trace.get_current_span()
    span.set_attribute("user.id", user_id)
    # Your logic here

```

## 2.12 Web Scraping with Scrapy

```

import scrapy

class ProductSpider(scrapy.Spider):
    name = 'products'
    start_urls = ['https://example.com/products']

    custom_settings = {
        'DOWNLOAD_DELAY': 2, # Respectful scraping
        'CONCURRENT_REQUESTS': 1
    }

    def parse(self, response):

```

```
for product in response.css('div.product'):
    yield {
        'name': product.css('h2::text').get(),
        'price': product.css('span.price::text').get(),
        'rating': product.css('div.rating::attr(data-rating)').get()
    }
```

## 2.13 Ethical Web Scraping

1. **Check robots.txt** - Respect website's scraping policy
2. **Rate Limiting** - Don't overwhelm servers
3. **User-Agent** - Identify yourself honestly
4. **Terms of Service** - Read and comply with ToS
5. **Personal Data** - Respect privacy laws (GDPR, CCPA)



### Warning

Scraping can violate ToS or copyright. Always consult legal counsel for commercial use.

## 2.14 Streaming Data with Kafka

```
from kafka import KafkaProducer, KafkaConsumer
import json

# Producer: Collect data
producer = KafkaProducer(
    bootstrap_servers=['localhost:9092'],
    value_serializer=lambda v: json.dumps(v).encode('utf-8')
)

event = {
    'user_id': '12345',
    'event_type': 'page_view',
    'page': '/products'
}
producer.send('user-events', event)

# Consumer: Process data
consumer = KafkaConsumer('user-events')
```

```
for message in consumer:  
    process_event(message.value)
```

## 2.15 Data Collection Best Practices

### 2.15.1 Schema Design with Pydantic

```
from pydantic import BaseModel, Field  
from datetime import datetime  
  
class UserEvent(BaseModel):  
    user_id: str  
    event_type: str  
    timestamp: datetime = Field(default_factory=datetime.utcnow)  
    properties: dict  
    session_id: Optional[str] = None
```

## 2.16 Sampling Strategies

```
import random  
  
class SamplingCollector:  
    def __init__(self, sample_rate=0.1):  
        self.sample_rate = sample_rate  
  
    def should_collect(self, event):  
        if random.random() < self.sample_rate:  
            return True  
        if event.get('priority') == 'high':  
            return True  
        return False
```

**When to use:** High-traffic systems where full collection is expensive

## 2.17 Privacy and Compliance

```

import hashlib

class PrivacyAwareCollector:
    def __init__(self, pii_fields=['email', 'phone']):
        self.pii_fields = pii_fields

    def anonymize(self, data):
        anonymized = data.copy()
        for field in self.pii_fields:
            if field in anonymized:
                value = str(anonymized[field])
                hashed = hashlib.sha256(value.encode()).hexdigest()
                anonymized[field] = hashed
        return anonymized

```

## 3 Part 2: Data Validation

Ensuring Data Quality and Reliability

### 3.1 Why Data Validation Matters

#### 3.1.1 The Cost of Bad Data

- Garbage In, Garbage Out
- Silent Failures
- Expensive Debugging
- Lost Trust

#### 3.1.2 Real-World Impact

**E-commerce:** Missing IDs → 30% drop in CTR

**Medical:** Wrong units → Misdiagnoses

**Fraud Detection:** Duplicates → 45% false positives

### 3.2 Types of Data Quality Issues

1. **Completeness** - Missing or null values
2. **Accuracy** - Incorrect or outdated values
3. **Consistency** - Contradictory data

4. **Timeliness** - Stale or outdated data
5. **Uniqueness** - Duplicate records
6. **Validity** - Violates business rules

### 3.3 Data Validation with Pydantic

```
from pydantic import BaseModel, Field, validator, EmailStr

class UserEvent(BaseModel):
    user_id: str = Field(..., min_length=1, max_length=100)
    email: EmailStr
    age: int = Field(..., ge=0, le=150)
    event_type: str

    @validator('event_type')
    def validate_event_type(cls, v):
        allowed = ['click', 'view', 'purchase', 'signup']
        if v not in allowed:
            raise ValueError(f'event_type must be one of {allowed}')
        return v

    @validator('timestamp')
    def timestamp_not_future(cls, v):
        if v > datetime.utcnow():
            raise ValueError('timestamp cannot be in the future')
        return v
```

### 3.4 Using Pydantic

```
try:
    event = UserEvent(
        user_id="user_123",
        email="user@example.com",
        age=25,
        event_type="click",
        timestamp=datetime.utcnow()
    )
    print(" Valid event")
```

```
except ValueError as e:  
    print(" Validation errors:", e.json())
```

### 3.5 Great Expectations Framework

```
import great_expectations as gx  
from great_expectations.dataset import PandasDataset  
  
df = pd.read_csv('user_events.csv')  
dataset = PandasDataset(df)  
  
# Define expectations  
dataset.expect_column_values_to_not_be_null('user_id')  
dataset.expect_column_values_to_be_unique('event_id')  
dataset.expect_column_values_to_be_in_set(  
    'event_type',  
    ['click', 'view', 'purchase'])  
dataset.expect_column_values_to_be_between('age', 0, 150)  
  
# Validate  
results = dataset.validate()
```

### 3.6 Pandera: Statistical Validation

```
import pandera as pa  
from pandera import Column, Check  
  
schema = pa.DataFrameSchema({  
    "user_id": Column(str, checks=[  
        Check.str_length(min_value=1, max_value=100)  
    ]),  
    "age": Column(int, checks=[  
        Check.in_range(min_value=0, max_value=150),  
        Check(lambda s: s.mean() > 18)  
    ]),  
    "purchase_amount": Column(float, checks=[  
        Check.greater_than_or_equal_to(0)
```

```

    ] , nullable=True)
})

validated_df = schema.validate(df)

```

### 3.7 Custom Validation Pipelines

```

from dataclasses import dataclass
from typing import List

@dataclass
class ValidationResult:
    passed: bool
    errors: List[str]
    warnings: List[str]

class DataValidationPipeline:
    def __init__(self):
        self.validators = []

    def add_validator(self, validator):
        self.validators.append(validator)
        return self

    def validate(self, df):
        all_errors = []
        for validator in self.validators:
            result = validator.validate(df)
            all_errors.extend(result.errors)
        return ValidationResult(
            passed=len(all_errors) == 0,
            errors=all_errors,
            warnings=[]
        )

```

### 3.8 Statistical Distribution Validation

```

class DistributionValidator:
    def __init__(self, reference_df, columns):
        self.reference_stats = {
            col: {
                'mean': reference_df[col].mean(),
                'std': reference_df[col].std()
            }
            for col in columns
        }

    def validate(self, df):
        errors = []
        for col, ref_stats in self.reference_stats.items():
            current_mean = df[col].mean()
            mean_shift = abs(current_mean - ref_stats['mean']) / ref_stats['std']

            if mean_shift > 2:
                errors.append(
                    f"{col}: Mean shifted by {mean_shift:.2f} std deviations"
                )
        return ValidationResult(passed=len(errors) == 0, errors=errors, warnings=[])


```

### 3.9 Data Quality Monitoring

```

import sentry_sdk
from datadog import statsd

class DataQualityMonitor:
    def monitor(self, df, dataset_name):
        result = self.pipeline.validate(df)

        # Send metrics
        statsd.gauge(f'{dataset_name}.row_count', len(df))
        statsd.gauge(f'{dataset_name}.validation.errors', len(result.errors))

        # Alert on failures
        if not result.passed:
            sentry_sdk.capture_message(
                f"Data validation failed for {dataset_name}",

```

```
        extras={"errors": result.errors}  
    )
```

## 4 Part 3: Data Labeling

Annotating Datasets with Ground Truth

### 4.1 What is Data Labeling?

**Data Labeling:** Adding meaningful tags or annotations to raw data

#### 4.1.1 Why It's Critical

- **Supervised Learning** requires labeled examples
- **Quality labels** → Better models
- **Consistent labels** → Reliable evaluation
- **Cost:** Often 60-80% of ML project time and budget

## 4.2 Types of Labeling Tasks

### 4.2.1 Classification

- Image: “cat” vs “dog”
- Text: “spam” vs “not spam”
- Audio: speaker identification

### 4.2.2 Object Detection

- Bounding boxes
- Keypoint annotation
- Polygon segmentation

### 4.2.3 Sequence Labeling

- Named Entity Recognition
- Part-of-speech tagging
- Time series anomalies

#### 4.2.4 Structured Prediction

- Semantic segmentation
- Dependency parsing
- Relationship extraction

### 4.3 Label Studio

**Label Studio:** Open-source, multi-modal annotation platform

#### 4.3.1 Key Features

- Multi-modal: Images, text, audio, video, time series
- Custom interfaces with XML config
- ML-assisted labeling and active learning
- Collaboration features
- Export: JSON, CSV, COCO, Pascal VOC, YOLO

#### 4.3.2 Installation

```
pip install label-studio  
label-studio start
```

### 4.4 Label Studio: Image Classification

```
<View>  
  <Image name="image" value="$image_url"/>  
  <Choices name="choice" toName="image">  
    <Choice value="Cat"/>  
    <Choice value="Dog"/>  
    <Choice value="Other"/>  
  </Choices>  
</View>
```

## 4.5 Label Studio: Object Detection

```
<View>
  <Image name="image" value="$image_url"
    zoom="true" zoomControl="true"/>
  <RectangleLabels name="label" toName="image">
    <Label value="Person" background="red"/>
    <Label value="Car" background="blue"/>
    <Label value="Bicycle" background="green"/>
  </RectangleLabels>
</View>
```

## 4.6 Label Studio: NER

```
<View>
  <Text name="text" value="$text"/>
  <Labels name="label" toName="text">
    <Label value="Person" background="red"/>
    <Label value="Organization" background="blue"/>
    <Label value="Location" background="green"/>
    <Label value="Date" background="orange"/>
  </Labels>
</View>
```

Example: “Apple Inc. CEO Tim Cook announced the new iPhone in Cupertino”

## 4.7 Alternative Labeling Tools

### 4.7.1 Labelbox

- Enterprise features
- Model-assisted labeling
- **Cons:** Expensive, vendor lock-in

### 4.7.2 CVAT

- Open-source by Intel
- Video annotation
- **Cons:** Setup complexity

#### **4.7.3 Prodigy**

- By spaCy creators
- Active learning built-in
- **Cons:** License cost

#### **4.7.4 Scale AI / SageMaker**

- Managed services
- Human workforce included
- **Cons:** Very expensive

### **4.8 Inter-Annotator Agreement**

**Problem:** Different annotators may label differently

**Solution:** Measure agreement to ensure quality

#### **4.8.1 Cohen's Kappa ( )**

Measures agreement between **two annotators** accounting for chance

$$\kappa = \frac{p_o - p_e}{1 - p_e}$$

**Interpretation:**

- 0.0-0.20: Slight
- 0.21-0.40: Fair
- 0.41-0.60: Moderate
- 0.61-0.80: Substantial
- 0.81-1.00: Almost perfect

### **4.9 Cohen's Kappa Implementation**

```

from sklearn.metrics import cohen_kappa_score

annotator1 = ['spam', 'ham', 'spam', 'ham', ...]
annotator2 = ['spam', 'ham', 'ham', 'ham', ...]

kappa = cohen_kappa_score(annotator1, annotator2)
print(f"Cohen's Kappa: {kappa:.3f}")

```

**Example Output:**

Cohen's Kappa: 0.745  
 → Substantial agreement

## 4.10 Fleiss' Kappa: Multiple Annotators

```

from statsmodels.stats.inter_rater import fleiss_kappa
import numpy as np

# Matrix: rows = items, columns = category counts
data = np.array([
    [3, 0], # Item 1: all 3 annotators said "ham"
    [2, 1], # Item 2: 2 said "ham", 1 said "spam"
    [0, 3], # Item 3: all 3 said "spam"
])

kappa = fleiss_kappa(data, method='fleiss')
print(f"Fleiss' Kappa: {kappa:.3f}")

```

## 4.11 Improving Agreement

### 4.11.1 1. Clear Guidelines

- Define each label precisely
- Provide examples and counter-examples
- Document edge cases
- Create decision trees for ambiguous cases

#### **4.11.2 2. Training Phase**

- Use gold standard datasets
- Review disagreements together
- Iterative calibration sessions

#### **4.11.3 3. Regular Consensus Meetings**

- Discuss disputed items
- Update guidelines based on patterns
- Track improvement over time

### **4.12 Sampling Strategies**

#### **4.12.1 Random Sampling**

```
import random
random.sample(dataset, n=100)
```

**Pros:** Unbiased **Cons:** May miss rare classes

#### **4.12.2 Stratified Sampling**

```
from sklearn.model_selection import train_test_split
_, sampled, _, _ = train_test_split(
    dataset, labels,
    train_size=len(dataset)-100,
    stratify=labels
)
```

**Pros:** Ensures class representation

### **4.13 Snowball Sampling**

```

from sklearn.neighbors import NearestNeighbors

class SnowballSampler:
    def sample(self, labeled_data, unlabeled_data, n=100):
        # Vectorize
        labeled_vectors = self.vectorizer.fit_transform(labeled_data)
        unlabeled_vectors = self.vectorizer.transform(unlabeled_data)

        # Find nearest neighbors
        nn = NearestNeighbors(n_neighbors=5)
        nn.fit(unlabeled_vectors)
        distances, indices = nn.kneighbors(labeled_vectors)

        return [unlabeled_data[i] for i in indices.flatten()[:n]]

```

**Use case:** Finding more examples of a rare class

## 4.14 Active Learning

**Active Learning:** Intelligently select which examples to label next

### 4.14.1 Process

1. Train model on small labeled set
2. Find **most informative** unlabeled examples
3. Label those examples
4. Retrain model
5. Repeat

**Goal:** Achieve good performance with minimal labeling

## 4.15 Uncertainty Sampling

```

class UncertaintySampler:
    def sample(self, X_unlabeled, n=10):
        probs = self.model.predict_proba(X_unlabeled)

        # Higher entropy = more uncertain
        uncertainties = -np.sum(probs * np.log(probs + 1e-10), axis=1)

```

```

# Select top-n most uncertain
uncertain_indices = np.argsort(uncertainties)[-n:]
return uncertain_indices

```

## 4.16 Active Learning Loop

```

def active_learning_loop(X_labeled, y_labeled, X_unlabeled, budget=100):
    model = MultinomialNB()
    sampler = UncertaintySampler(model)

    for iteration in range(budget // 10):
        # Train model
        model.fit(X_labeled, y_labeled)

        # Select informative examples
        indices = sampler.sample(X_unlabeled, n=10)

        # Get human labels
        new_X = [X_unlabeled[i] for i in indices]
        new_y = get_human_labels(new_X)

        # Update datasets
        X_labeled.extend(new_X)
        y_labeled.extend(new_y)

        # Evaluate
        accuracy = evaluate_model(model)
        print(f"Iteration {iteration}: Accuracy = {accuracy:.3f}")

```

## 4.17 Active Learning Benefits

### 4.17.1 Cost Reduction

Random Sampling: 10,000 labels → 85% accuracy

Active Learning: 2,000 labels → 85% accuracy

**Savings: 80% fewer labels!**

#### 4.17.2 Real-world Savings

- 10k labels  $\times \$0.50 = \$5,000$
- 2k labels  $\times \$0.50 = \$1,000$
- Saved: **\$4,000**

**Key Insight:** Active learning reaches target accuracy with 5-10 $\times$  fewer labels

### 4.18 Weak Supervision with Snorkel

**Weak Supervision:** Use noisy/heuristic labels instead of manual annotation

```
from snorkel.labeling import labeling_function

SPAM = 1
HAM = 0
ABSTAIN = -1

@labeling_function()
def lf_contains_money(x):
    money_keywords = ['$', 'money', 'cash', 'prize']
    return SPAM if any(kw in x.text.lower() for kw in money_keywords) else ABSTAIN

@labeling_function()
def lf_short_message(x):
    return SPAM if len(x.text.split()) < 10 else ABSTAIN

@labeling_function()
def lf_known_sender(x):
    trusted_domains = ['company.com', 'university.edu']
    return HAM if any(d in x.sender for d in trusted_domains) else ABSTAIN
```

### 4.19 Combining Weak Labels

```
from snorkel.labeling.model import LabelModel

# Apply labeling functions
lfs = [lf_contains_money, lf_short_message, lf_known_sender]
L_train = applier.apply(df=df)
```

```

# Combine noisy labels
label_model = LabelModel(cardinality=2)
label_model.fit(L_train=L_train, n_epochs=500)

# Get predictions
predicted_labels = label_model.predict(L=L_train)

```

**Advantages:** Fast, Flexible, No manual labeling needed

## 4.20 Pre-labeling with Models

```

class PreLabelingPipeline:
    def pre_label(self, unlabeled_data, confidence_threshold=0.9):
        predictions = self.model.predict_proba(unlabeled_data)

        auto_labeled = []
        needs_review = []

        for i, probs in enumerate(predictions):
            max_prob = max(probs)

            if max_prob >= confidence_threshold:
                auto_labeled.append({
                    'data': unlabeled_data[i],
                    'label': np.argmax(probs),
                    'source': 'model'
                })
            else:
                needs_review.append(unlabeled_data[i])

        return auto_labeled, needs_review

```

## 5 Part 4: Data Augmentation

Expanding Training Datasets Strategically

## 5.1 Why Data Augmentation?

### 5.1.1 The Problem

- Limited labeled data is expensive
- Class imbalance leads to biased models
- Overfitting on small datasets
- Rare events underrepresented

### 5.1.2 The Solution

**Data Augmentation:** Creating new training examples by applying transformations

**Benefits:** Increase dataset size, Improve generalization, Reduce overfitting, Balance classes

## 5.2 Image Data Augmentation

```
from torchvision import transforms

transform = transforms.Compose([
    # Geometric
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomRotation(degrees=15),
    transforms.RandomResizedCrop(size=224, scale=(0.8, 1.0)),

    # Color
    transforms.ColorJitter(
        brightness=0.2, contrast=0.2,
        saturation=0.2, hue=0.1
    ),

    # Blurring
    transforms.GaussianBlur(kernel_size=3),

    # Normalization
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225])
])
```

### 5.3 Advanced: Albumentations

```
import albumentations as A

transform = A.Compose([
    A.RandomResizedCrop(height=224, width=224),
    A.HorizontalFlip(p=0.5),
    A.ShiftScaleRotate(shift_limit=0.0625, scale_limit=0.1, rotate_limit=15),
    A.ElasticTransform(p=0.3),

    # Color
    A.RandomBrightnessContrast(p=0.5),
    A.HueSaturationValue(p=0.5),

    # Blur and noise
    A.OneOf([
        A.MotionBlur(p=0.5),
        A.GaussianBlur(p=0.5),
    ], p=0.3),
    A.GaussNoise(p=0.3),

    # Weather
    A.RandomRain(p=0.2),
    A.RandomFog(p=0.2),
])
])
```

### 5.4 Cutout, Mixup, CutMix

#### 5.4.1 Cutout

Random rectangular masks

```
A.CoarseDropout(
    max_holes=8,
    max_height=32,
    p=0.5
)
```

### 5.4.2 Mixup

Blend two images

```
def mixup(x1, y1, x2, y2):
    lam = np.random.beta(0.2, 0.2)
    x = lam * x1 + (1 - lam) * x2
    y = lam * y1 + (1 - lam) * y2
    return x, y
```

### 5.4.3 CutMix

Paste regions

```
def cutmix(x1, y1, x2, y2):
    # Cut and paste
    # region from x2 to x1
    ...
```

## 5.5 Text Data Augmentation

### 5.5.1 Synonym Replacement

```
from nltk.corpus import wordnet

def get_synonyms(word):
    synonyms = set()
    for syn in wordnet.synsets(word):
        for lemma in syn.lemmas():
            synonyms.add(lemma.name())
    return list(synonyms)

def synonym_replacement(text, n=2):
    words = text.split()
    random_indices = random.sample(range(len(words)), n)

    for idx in random_indices:
        synonyms = get_synonyms(words[idx])
        if synonyms:
            words[idx] = random.choice(synonyms)
```

```
    return ' '.join(words)
```

## 5.6 EDA: Easy Data Augmentation

```
class EDA:  
    @staticmethod  
    def random_insertion(words, n=1):  
        """Insert n random synonyms"""  
        ...  
  
    @staticmethod  
    def random_swap(words, n=1):  
        """Swap n pairs of words"""  
        ...  
  
    @staticmethod  
    def random_deletion(words, p=0.1):  
        """Delete each word with probability p"""  
        return [w for w in words if random.random() > p]
```

## 5.7 Back-Translation

```
from transformers import MarianMTModel, MarianTokenizer  
  
class BackTranslator:  
    def __init__(self, intermediate_lang='fr'):  
        # English → French  
        self.forward_model = MarianMTModel.from_pretrained(  
            f'Helsinki-NLP/opus-mt-en-{intermediate_lang}'  
        )  
        # French → English  
        self.backward_model = MarianMTModel.from_pretrained(  
            f'Helsinki-NLP/opus-mt-{intermediate_lang}-en'  
        )  
  
    def augment(self, text):  
        # Translate to French and back
```

```

        intermediate = self.forward_model.translate(text)
        return self.backward_model.translate(intermediate)

```

Result: “The weather is beautiful” → “The weather is nice”

## 5.8 Contextual Augmentation with BERT

```

from transformers import pipeline

class ContextualAugmentor:
    def __init__(self):
        self.unmasker = pipeline('fill-mask', model='bert-base-uncased')

    def augment(self, text):
        words = text.split()
        mask_idx = random.randint(0, len(words) - 1)

        masked_words = words.copy()
        masked_words[mask_idx] = '[MASK]'

        predictions = self.unmasker(' '.join(masked_words))

        return [pred['sequence'] for pred in predictions[:3]]

```

## 5.9 Audio Augmentation

```

import librosa

class AudioAugmentor:
    @staticmethod
    def add_noise(audio, noise_factor=0.005):
        noise = np.random.randn(len(audio))
        return audio + noise_factor * noise

    @staticmethod
    def time_stretch(audio, rate=1.0):
        return librosa.effects.time_stretch(audio, rate=rate)

```

```

@staticmethod
def pitch_shift(audio, sr=22050, n_steps=2):
    return librosa.effects.pitch_shift(audio, sr=sr, n_steps=n_steps)

```

## 5.10 Time Series Augmentation

```

class TimeSeriesAugmentor:
    @staticmethod
    def jittering(x, sigma=0.03):
        """Add Gaussian noise"""
        return x + np.random.normal(0, sigma, x.shape)

    @staticmethod
    def scaling(x, sigma=0.1):
        """Multiply by random factor"""
        factor = np.random.normal(1, sigma, size=(x.shape[0], 1))
        return x * factor

    @staticmethod
    def time_warping(x, sigma=0.2):
        """Warp time dimension"""
        # Use cubic spline interpolation
        ...

```

## 5.11 SMOTE for Tabular Data

**SMOTE:** Synthetic Minority Over-sampling Technique

```

from imblearn.over_sampling import SMOTE

# Imbalanced dataset: {0: 900, 1: 100}
smote = SMOTE(sampling_strategy='auto', random_state=42)
X_resampled, y_resampled = smote.fit_resample(X_train, y_train)

# Now balanced: {0: 900, 1: 900}

```

**How it works:**

1. For each minority sample, find k nearest neighbors

2. Randomly select one neighbor
3. Create synthetic sample along line segment

## 5.12 SMOTE Variants

### 5.12.1 ADASYN

Adaptive Synthetic Sampling

```
from imblearn.over_sampling import ADASYN

adasyn = ADASYN(random_state=42)
X_res, y_res = adasyn.fit_resample(X, y)
```

**Advantage:** Generates more samples in harder-to-learn regions

### 5.12.2 BorderlineSMOTE

Focus on decision boundary

```
from imblearn.over_sampling import BorderlineSMOTE

bsmote = BorderlineSMOTE(random_state=42)
X_res, y_res = bsmote.fit_resample(X, y)
```

**Advantage:** More conservative, focuses on boundary

## 5.13 Generative Models

```
from diffusers import StableDiffusionPipeline

pipe = StableDiffusionPipeline.from_pretrained(
    "stabilityai/stable-diffusion-2-1"
)

def generate_augmented_images(prompt, n=10):
    images = []
    for i in range(n):
        image = pipe(prompt, num_inference_steps=50).images[0]
```

```

        images.append(image)
    return images

# Generate synthetic training data
synthetic = generate_augmented_images(
    "A high-quality photo of a defective product",
    n=100
)

```

**Use cases:** Rare defects, medical imaging, artistic styles

## 5.14 Best Practices

### 5.14.1 1. Domain-Appropriate

**Good:** Horizontal flip for general images    **Bad:** Horizontal flip for text/digits (changes meaning)

### 5.14.2 2. Preserve Semantics

**Bad:** Extreme rotation for digits ( $6 \rightarrow 9$ )    **Good:** Slight rotation ( $\pm 15^\circ$ )

### 5.14.3 3. Validation Split First

```

#   Correct
train, val = train_test_split(data)
aug_train = augment(train) # Only augment training

#   Wrong (data leakage!)
aug_data = augment(data)
train, val = train_test_split(aug_data)

```

## 5.15 Monitor Augmentation Impact

```

def evaluate_augmentation(model, train_data, val_data, aug_levels):
    results = []

    for strength in aug_levels:
        aug_train = augment(train_data, strength=strength)
        model.fit(aug_train)
        val_acc = model.evaluate(val_data)
        results.append({'strength': strength, 'accuracy': val_acc})

    # Plot results to find optimal augmentation
    plt.plot([r['strength'] for r in results],
              [r['accuracy'] for r in results])
    plt.xlabel('Augmentation Strength')
    plt.ylabel('Validation Accuracy')

```

## 5.16 Class-Specific Augmentation

```

class ClassBalancedAugmentor:
    def __init__(self, target_samples_per_class=1000):
        self.target = target_samples_per_class

    def augment(self, X, y):
        X_aug, y_aug = [], []

        for cls in np.unique(y):
            X_cls = X[y == cls]
            current_count = len(X_cls)

            if current_count < self.target:
                n_to_generate = self.target - current_count
                augmented = self.generate_samples(X_cls, n_to_generate)
                X_aug.extend(augmented)
                y_aug.extend([cls] * len(augmented))

        return np.array(X_aug), np.array(y_aug)

```

## 6 Summary

### 6.1 Key Takeaways

#### 6.1.1 Data Collection

- Instrument systems for automatic capture
- Use appropriate tools (analytics, APM, CDC)
- Implement buffering, batching, error handling
- Respect privacy and comply with regulations

#### 6.1.2 Data Validation

- Validate early and often
- Use schema validation (Pydantic, Pandera)
- Monitor data quality metrics
- Set up alerts for anomalies

### 6.2 Key Takeaways (cont.)

#### 6.2.1 Data Labeling

- Use appropriate tools (Label Studio, CVAT)
- Measure inter-annotator agreement (Cohen's Kappa)
- Apply smart sampling (active learning)
- Leverage weak supervision when possible

#### 6.2.2 Data Augmentation

- Choose domain-appropriate transformations
- Preserve label semantics
- Augment only training data
- Monitor impact on performance

! Important

High-quality data is the foundation of successful AI systems. Invest time in getting it right!

## 6.3 Hands-On Lab

### 6.3.1 Complete Data Pipeline

**Tasks:**

1. Data Collection (30 min) - Scrape reviews, set up logging
2. Data Validation (30 min) - Create schemas, build validation pipeline
3. Data Labeling (45 min) - Label Studio, calculate agreement, active learning
4. Data Augmentation (30 min) - Apply EDA, back-translation, measure impact

**Deliverable:** Jupyter notebook with complete pipeline

## 6.4 Additional Resources

### 6.4.1 Documentation

- Label Studio: <https://labelstud.io/guide/>
- Great Expectations: <https://docs.greatexpectations.io/>
- Albulmentations: <https://albulmentations.ai/docs/>
- Snorkel: <https://www.snorkel.org/>

### 6.4.2 Key Papers

- “SMOTE: Synthetic Minority Over-sampling Technique” (2002)
- “Snorkel: Rapid Training Data Creation with Weak Supervision” (2017)
- “Active Learning Literature Survey” (2009)

## 7 Questions?

### 7.1 Thank You!

**Next:** Reproducibility & Versioning

---

Contact: [nipun.batra@iitgn.ac.in](mailto:nipun.batra@iitgn.ac.in)