

Building ML APIs with FastAPI

Week 10 · CS 203: Software Tools and Techniques for AI

Prof. Nipun Batra
IIT Gandhinagar

The Problem: Sharing Your Model

You've built an amazing movie predictor:

- Trained a model that predicts Netflix success
- Works great in your Jupyter notebook

But now:

- Your professor wants to try it
- A mobile app developer wants to use it
- The marketing team wants a dashboard

Challenge: They can't run your notebook!

The Solution: APIs

API = Application Programming Interface

Think of it like a restaurant:

- You (client) order food through a menu (API)
- Kitchen (server) prepares the food
- You don't need to know HOW the kitchen works

For ML:

Request	Response
Send movie features	Get prediction back
<code>{"budget": 100, "genre": "Action"}</code>	<code>{"success": true, "confidence": 0.85}</code>

Why APIs Are Revolutionary

An API is a universal translator. It doesn't matter if your model is in Python, if your client is a mobile app written in Swift, or if your dashboard uses JavaScript. Everyone speaks HTTP. Everyone understands JSON. Your ML model becomes accessible to the entire world.

Python Model ↔ API ↔ Mobile App (Swift)



Web App (JavaScript)



Another Service (Go)

One model, infinite consumers.

HTTP: The Language of the Web

Every web request has:

Method	Purpose	Example
GET	Read data	Get user profile
POST	Send data	Submit prediction
PUT	Update data	Update settings
DELETE	Remove data	Delete account

For ML predictions: We mostly use **POST** (send data, get prediction)

What is FastAPI?

FastAPI = Python framework for building APIs

Why FastAPI?

1. **Fast**: High performance (like Node.js)
2. **Easy**: Write Python, get web APIs
3. **Auto-docs**: Interactive documentation for free
4. **Validation**: Automatic input checking

```
pip install "fastapi[standard]"
```

Your First API: Hello World

Create `app.py` :

```
from fastapi import FastAPI

app = FastAPI()

@app.get("/")
def hello():
    return {"message": "Hello, World!"}
```

Run it:

```
fastapi dev app.py
```

Visit: <http://localhost:8000>

Understanding the Code

```
from fastapi import FastAPI

app = FastAPI() # Create the app

@app.get("/")    # When someone visits "/"
def hello():     # Run this function
    return {"message": "Hello, World!"} # Return JSON
```

The `@app.get("/")` decorator:

- `@app.get` = Handle GET requests
- `"/"` = At the root URL (<http://localhost:8000/>)

Auto-Generated Documentation

Visit: <http://localhost:8000/docs>

You get **Swagger UI** for free:

- See all your endpoints
- Test them interactively
- View request/response schemas

Try it: Click "Try it out" and execute!

Path Parameters

Dynamic URLs:

```
@app.get("/movies/{movie_id}")
def get_movie(movie_id: int):
    return {"movie_id": movie_id, "title": "Movie " + str(movie_id)}
```

Examples:

- `/movies/42` → `{"movie_id": 42, "title": "Movie 42"}`
- `/movies/123` → `{"movie_id": 123, "title": "Movie 123"}`

Note: FastAPI automatically converts `"42"` to integer `42`

Query Parameters

Optional parameters in URL:

```
@app.get("/movies")
def search_movies(genre: str = None, limit: int = 10):
    return {
        "genre": genre,
        "limit": limit,
        "results": ["Movie 1", "Movie 2"]
    }
```

Examples:

- `/movies` → Uses defaults
- `/movies?genre=Action` → Filter by genre
- `/movies?genre=Comedy&limit=5` → Both parameters

POST Requests: Sending Data

For predictions, we need to send data:

```
from pydantic import BaseModel

class MovieInput(BaseModel):
    genre: str
    budget: float
    runtime: int

@app.post("/predict")
def predict(movie: MovieInput):
    # Prediction logic here
    return {"success": True, "confidence": 0.85}
```

Client sends JSON:

```
{"genre": "Action", "budget": 150.0, "runtime": 120}
```

What is Pydantic?

Pydantic = Data validation library

```
from pydantic import BaseModel

class MovieInput(BaseModel):
    genre: str      # Must be a string
    budget: float   # Must be a number
    runtime: int    # Must be an integer
```

FastAPI uses Pydantic to:

1. Validate incoming data
2. Show schema in docs
3. Return clear error messages

Pydantic: Your API's Bouncer

Pydantic is like a bouncer at a club with a strict dress code. No ID? Go away. Wrong type? Go away. Budget is negative? Go away. This validation happens BEFORE your code runs, so you never have to write `if budget < 0: return error`. Pydantic does it automatically. Your model only sees clean, validated data.

Raw Request JSON

```
{"budget": "abc", ...} —→  
✗  
422 Error (never reaches your code)  
{"budget": -10, ...} —→  
✗  
422 Error (never reaches your code)  
{"budget": 100, ...} —→ ✓ MovieInput → predict(movie)
```

Pydantic

Your Code

Trust your inputs. Pydantic already checked them.

Pydantic Field Validation

Add constraints:

```
from pydantic import BaseModel, Field

class MovieInput(BaseModel):
    genre: str
    budget: float = Field(gt=0, lt=500) # Between 0-500
    runtime: int = Field(ge=60, le=240) # 60-240 minutes
    is_sequel: bool = False # Optional with default
```

If budget is -10:

```
{"detail": "budget must be greater than 0"}
```

Serving Your ML Model

The pattern:

```
import joblib
from fastapi import FastAPI

app = FastAPI()
model = None # Global variable

@app.on_event("startup")
def load_model():
    global model
    model = joblib.load("movie_model.pkl")
    print("Model loaded!")
```

Why load at startup?

- Load once, use many times
- Don't reload for each request

Complete Prediction Endpoint

```
@app.post("/predict")
def predict(movie: MovieInput):
    # Prepare features
    features = [
        movie.budget,
        movie.runtime,
        1 if movie.is_sequel else 0
    ]

    # Make prediction
    prediction = model.predict([features])[0]
    probability = model.predict_proba([features])[0]

    return {
        "prediction": "Success" if prediction == 1 else "Risky",
        "confidence": float(max(probability))
    }
```

Complete Example: Movie Predictor API

```
from fastapi import FastAPI, HTTPException
from pydantic import BaseModel, Field
import joblib

app = FastAPI(title="Movie Success Predictor")
model = None

class MovieInput(BaseModel):
    budget: float = Field(gt=0, description="Budget in millions")
    runtime: int = Field(ge=60, le=240, description="Runtime in minutes")
    is_sequel: bool = False

class PredictionOutput(BaseModel):
    prediction: str
    confidence: float

@app.on_event("startup")
def load_model():
    global model
    model = joblib.load("movie_model.pkl")
```

Complete Example (continued)

```
@app.get("/health")
def health_check():
    return {
        "status": "healthy",
        "model_loaded": model is not None
    }

@app.post("/predict", response_model=PredictionOutput)
def predict(movie: MovieInput):
    if model is None:
        raise HTTPException(status_code=503, detail="Model not loaded")

    features = [[movie.budget, movie.runtime, int(movie.is_sequel)]]
    prediction = model.predict(features)[0]
    probability = model.predict_proba(features)[0]

    return {
        "prediction": "Success" if prediction == 1 else "Risky",
        "confidence": float(max(probability))
    }
```

Error Handling

What if something goes wrong?

```
from fastapi import HTTPException

@app.post("/predict")
def predict(movie: MovieInput):
    if model is None:
        raise HTTPException(
            status_code=503,
            detail="Model not available"
        )

    try:
        result = model.predict(...)
        return {"prediction": result}
    except Exception as e:
        raise HTTPException(
            status_code=500,
            detail=f"Prediction failed: {str(e)}"
        )
```

HTTP Status Codes

Code	Meaning	When to use
200	OK	Request succeeded
400	Bad Request	Invalid input
404	Not Found	Resource doesn't exist
422	Validation Error	Pydantic validation failed
500	Server Error	Something crashed
503	Unavailable	Model not loaded

Status Codes: The Body Language of APIs

Status codes are how APIs communicate without words. Just like you can tell if someone is happy or angry from their expression, status codes tell you instantly if a request worked or failed - before you even read the response body. Learn the families: 2xx = success, 4xx = your fault, 5xx = server's fault.

Client: "Give me prediction for budget=-50"

Server: 422



(Validation Error - your fault!)

Client: "Give me prediction for budget=100"

Server: 500



(Server Error - our fault, sorry!)

Client: "Give me prediction for budget=100"

Server: 200



(Success!)

Health Check Endpoint

Always add a health check:

```
@app.get("/health")
def health_check():
    return {
        "status": "healthy",
        "model_loaded": model is not None,
        "version": "1.0.0"
    }
```

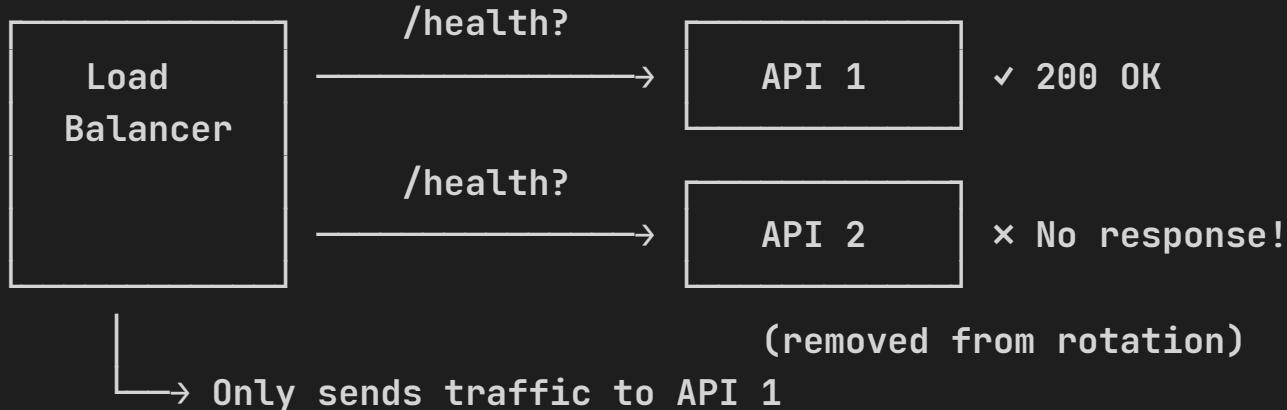
Why?

- Kubernetes/Docker checks if app is running
- Load balancers know where to send traffic
- Monitoring tools track uptime

The Heartbeat Analogy

A **health check** is your API's **pulse**. In production, automated systems constantly ping your `/health` endpoint - like a doctor checking your heartbeat. No response? The system knows something is wrong and can restart your service or redirect traffic. Without this, your API could be dead and nobody would know until users complain.

Load Balancer pings `/health` every 30 seconds:



Testing Your API

Use FastAPI's TestClient:

```
from fastapi.testclient import TestClient
from app import app

client = TestClient(app)

def test_health():
    response = client.get("/health")
    assert response.status_code == 200
    assert response.json()["status"] == "healthy"

def test_predict():
    response = client.post("/predict", json={
        "budget": 100, "runtime": 120, "is_sequel": False
    })
    assert response.status_code == 200
    assert "prediction" in response.json()
```

Running Tests

Install pytest:

```
pip install pytest
```

Create `test_app.py` with your tests

Run:

```
pytest test_app.py -v
```

Output:

```
test_app.py::test_health PASSED
test_app.py::test_predict PASSED
```

Running in Production

Development server (for testing):

```
fastapi dev app.py # Auto-reload on changes
```

Production server:

```
uvicorn app:app --host 0.0.0.0 --port 8000
```

With multiple workers:

```
uvicorn app:app --host 0.0.0.0 --port 8000 --workers 4
```

Docker Deployment

Create Dockerfile :

```
FROM python:3.10-slim

WORKDIR /app

COPY requirements.txt .
RUN pip install -r requirements.txt

COPY . .

CMD ["uvicorn", "app:app", "--host", "0.0.0.0", "--port", "8000"]
```

Build and run:

```
docker build -t movie-api .
docker run -p 8000:8000 movie-api
```

Calling Your API from Python

Using requests library:

```
import requests

# Make prediction
response = requests.post(
    "http://localhost:8000/predict",
    json={
        "budget": 100,
        "runtime": 120,
        "is_sequel": False
    }
)

print(response.json())
# {"prediction": "Success", "confidence": 0.85}
```

Calling Your API from JavaScript

```
// In a web app
fetch("http://localhost:8000/predict", {
  method: "POST",
  headers: {"Content-Type": "application/json"},
  body: JSON.stringify({
    budget: 100,
    runtime: 120,
    is_sequel: false
  })
})
.then(response => response.json())
.then(data => console.log(data));
```

CORS: Allowing Web Apps to Call Your API

Problem: Browsers block requests to different domains

Solution: Enable CORS

```
from fastapi.middleware.cors import CORSMiddleware

app.add_middleware(
    CORSMiddleware,
    allow_origins=["*"], # Allow all (use specific domains in production)
    allow_methods=["*"],
    allow_headers=["*"],
)
```

API Best Practices

1. Always validate inputs

```
budget: float = Field(gt=0, lt=500)
```

2. Return consistent responses

```
{"prediction": ..., "confidence": ...}
```

3. Add health checks

```
@app.get("/health")
```

4. Handle errors gracefully

```
except Exception as e:  
    raise HTTPException(500, detail=str(e))
```

5. Document your API (FastAPI does this automatically!)

Summary

Concept	What it does
FastAPI	Python framework for APIs
Pydantic	Data validation
GET	Read data
POST	Send data (predictions)
HTTPException	Handle errors
/health	Check if API is running
TestClient	Test your API

Lab Preview

This week you'll:

1. Create a FastAPI app from scratch
2. Add input validation with Pydantic
3. Serve your Netflix movie predictor model
4. Add error handling and health checks
5. Write tests for your API
6. Deploy with Docker (optional)

Result: A working ML API you can share with anyone!

Questions?

Key takeaways:

- FastAPI makes building ML APIs easy
- Pydantic validates your inputs automatically
- Always add health checks and error handling
- Test your API before deployment

Next week: Git and CI/CD